

Practical Machine Learning

Erik Rehnberg Steeb

8/15/2020

Setup R session

```
## Loading required package: lattice

## Loading required package: ggplot2

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
##
##     margin

## corrplot 0.84 loaded

## Loading required package: foreach

## Loading required package: iterators

## Loading required package: parallel
```

Load and Clean Data

First step is to load and clean the data

```
testData <- read.csv("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv")
trainData <- read.csv("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv")
# Remove columns with NAs
testData <- testData[, colSums(is.na(testData)) == 0]
trainData <- trainData[, colSums(is.na(trainData)) == 0]
# Remove columns that almost certainly do not have anything to do with prediction
testRemNames <- colnames(testData[,1:7])
trainRemNames <- colnames(trainData[,1:7])
testData <- testData[,!(names(testData) %in% testRemNames)]
```

```
trainData <- trainData [,!(names(trainData) %in% trainRemNames)]
# Remove empty variables from trainData
trainClasse <- trainData$classe
trainData <- trainData[sapply(trainData, is.numeric)]
trainData$classe <- trainClasse
trainData$classe <- as.factor(trainData$classe)
```

Slice Data

Slice data into training and validation test sets

```
set.seed(221989) # For reproducible purposes
inTrain <- createDataPartition(trainData$classe, p=0.70, list=F)
trueTrain <- trainData[inTrain,]
validateData <- trainData[-inTrain,]
```

Create the Model

I'll be using the Random Forest method, since it selects important variables itself (and is the default for Caret).

```
trainctrl <- trainControl(verboseIter = FALSE)
rfFit <- train(x = trueTrain[,1:52], y = trueTrain$classe,
              method = "rf",
              trControl = trainctrl
)

rfFit
```

```
## Random Forest
##
## 13737 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 13737, 13737, 13737, 13737, 13737, 13737, ...
## Resampling results across tuning parameters:
##
##  mtry  Accuracy  Kappa
##    2    0.9881654 0.9850270
##   27    0.9874988 0.9841830
##   52    0.9804433 0.9752563
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

Cross Validate

Now that we have a model, I'll estimate the accuracy on our validation dataset.

```

validateData$classe <- as.factor(validateData$classe)

predict <- predict(rfFit, validateData)
cMatrix <- confusionMatrix(validateData$classe, predict)

cMatrix

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A    B    C    D    E
##           A 1673    0    0    0    1
##           B   10 1119   10    0    0
##           C    0    6 1019    1    0
##           D    0    0   24  939    1
##           E    0    0    0    1 1081
##
## Overall Statistics
##
##           Accuracy : 0.9908
##           95% CI : (0.988, 0.9931)
##           No Information Rate : 0.286
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9884
##
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity          0.9941   0.9947   0.9677   0.9979   0.9982
## Specificity          0.9998   0.9958   0.9986   0.9949   0.9998
## Pos Pred Value       0.9994   0.9824   0.9932   0.9741   0.9991
## Neg Pred Value       0.9976   0.9987   0.9930   0.9996   0.9996
## Prevalence           0.2860   0.1912   0.1789   0.1599   0.1840
## Detection Rate       0.2843   0.1901   0.1732   0.1596   0.1837
## Detection Prevalence 0.2845   0.1935   0.1743   0.1638   0.1839
## Balanced Accuracy     0.9969   0.9952   0.9831   0.9964   0.9990

```

This gives us a very low estimated out-of-sample error rate - something like 0.03% of instances are likely to be misclassified.

Test the model

Time to actually test the model using the testData set

```

# Remove last column of testData
testData1 <- testData[,1:52]
predictTest <- predict(rfFit, testData)

predictTest


```

```
## [1] B A B A A E D B A A B C B A E E A B B B  
## Levels: A B C D E
```

Appendix: Figures

1. Correlation Matrix Visualization

```
corrPlot <- cor(trainData[, -length(names(trainData))])  
corrplot(corrPlot, method="color", tl.pos='n')
```



2. Decision Tree Visualization

```
treeModel <- rpart(classe ~ ., data=trainData, method="class")  
prp(treeModel) # fast plot
```

